

Where Next? Predicting the Next Unaffordable U.S. City Using Housing & Economic Indicators

Author: Brandon Smith

Course: DATA 510 - Capstone Project

Date: August 14th, 2026

Abstract

This project develops a housing-affordability risk framework to identify metro areas with patterns linked to affordability stress and potential future deterioration. This was completed using training data from three cities that have already experienced such collapse: Austin, TX; Boise, ID; and Tampa, FL. The project engineers temporal housing and economic features, trains an XGBoost binary classifier, and applies the resulting model to score 371 additional metro areas. SHAP values explain feature importance. Results show that the 5-year price-to-income change is the strongest feature associated with concurrent unaffordability, followed by affordability momentum and HPI year-over-year growth. The explanatory model (Model 1) serves as the project's primary answer to the research question. Model 1's dominant feature, the 5-year change in price-to-income ratio, comes from the same underlying PTI series used to construct the target label; as a result, Model 1 is best understood as classifying concurrent unaffordability rather than forecasting future collapse with meaningful lead time. This distinction is addressed directly in the Discussion and framed as a priority for future work.

Abstract	1
1. Introduction & Problem Statement	4
Primary Research Question	4
Data Scheme	4
Reproducibility	5
2. Literature Review & Context	5
The Affordability Crisis in Context	5
HPI as a Leading Indicator	5
Momentum-Based Detection	5
3. Data Sources & Warehousing	5
4. Affordability Collapse Definition & Threshold Selection	6
Condition 1 - Price Acceleration	6
Condition 2 - Sustained Momentum	7
Condition 3 - Affordability Burden (Primary Label Criterion)	7
Table 1: Affordability Collapse Onset Dates by PTI Threshold	7
5. Exploratory Data Analysis	7
Figure 1: Price-to-Income Ratio Trajectories (2010-2024) - Austin, Boise, Tampa	7
YoY ZHVI Growth Trends	8
Population Velocity	8
EDA-To-Feature Engineering Bridge	8
6. Feature Engineering	9
Features Engineered	9
Table 2: Feature Coverage by City	10
7. Predictive Modeling	10
Model 1 - Explanatory model (All 3 Training Cities, 4 Features)	10
Table 3: Model 1 Performance (Validation Set)	11
SHAP Feature Importance - Model 1	11
Figure 2: SHAP Summary Plot - Model 1	11
Table 4: Mean Absolute SHAP Values - Model 1	12
Cross-Table Synthesis:	12
Model 2 - Generalization/Scoring Model (Austin & Boise, 3 Features)	13
Table 5: Model 2 Performance (Validation Set)	13
Figure 3: SHAP Summary Plot - Model 2	14
Holdout City Risk Scoring	15
Figure 4: Holdout City Risk Scores	15
Why Bozeman, MT Ranks First in the City Risk Index	16
What a False Positive or False Negative Means for Stakeholders	16

8. Ethics & Limitations	17
Ethical Considerations	17
Limitations:	17
9. Conclusion	18
References	20
Appendix	22
Appendix A: Key Variables Definitions	22
Appendix B: Model Architecture Summary	22

1. Introduction & Problem Statement

Housing affordability in American cities has deteriorated sharply over the past decade. Cities like Austin, Boise, and Tampa were once affordable, mid-sized metros. Yet they underwent a rapid multi-year affordability collapse driven by population influx, constrained housing supply, and surging home values. Working- and middle-class residents were displaced as housing costs outpaced income growth.

Despite the severity of this trend, no reliable early-warning system exists to flag which cities are *next* to tip into unaffordability before displacement is already underway. Policy interventions like zoning reform, housing vouchers, and inclusionary requirements require years of lead time. City planners, regional lenders, and housing advocates are currently operating reactively, not proactively.

Primary Research Question

Which socioeconomic and housing market indicators best predict the onset of affordability collapse in mid-sized U.S. metros, and can those indicators identify at-risk cities before displacement begins?

Data Scheme

To allow peers to trace this project's logic from raw source CSVs to final outputs without reconstructing it from the notebook cells, the pipeline is summarized below:

- Merged dataset: 83,699 rows x 25 columns, quarterly observations across all U.S. CBSAs, 2010-2024, stored at data/final/merged_full.csv
- Collapse definition: `is_unaffordable = price_to_income_ratio >= 5.0`; `collapse_onset = first quarter is_unaffordable flips from False to Positive after a period of False`.
- Train/test logic (Model 1): The final reported Model 1 results use a stratified 70/30 random train/test split across all eligible observations from Austin, Boise, and Tampa, with `random_state = 42`. Stratification preserves the proportion of affordable and unaffordable observations in both sets. We created an onset-based split during initial data preparation but did not use it to generate the results reported in Table 3.
- Train/test logic (Model 2): The final reported Model 2 results use a 70/30 random train/test split across eligible observations from Austin and Boise, with `random_state = 42`. Stratification preserves the class distribution of the `is_unaffordable` target.
- Holdout set: All 371 non-training CBSAs with completed data for the 3 sourcing features (`hpi_yoy_lag`, `pop_velocity_lag`, `pop_acceleration_lag`), no grounded-truth labels available

Reproducibility

All data pipeline scripts, feature engineering code, model-training notebooks, figures, and output tables are available at: [GITHUB](#). The final models use `random_state = 42` to support reproducibility.

2. Literature Review & Context

The Affordability Crisis in Context

Research by the Urban Institute, HUD, and Zillow consistently identifies the price-to-income ratio (PTI) as one of the most reliable cross-market metrics for housing stress. A household spending more than 30% of gross income on housing is considered “cost-burdened” (HUD standard); a home priced at more than 3x median household income is historically considered unaffordable by conventional mortgage underwriting standards.

The 5.0x PTI threshold used in this project aligns with academic and industry literature showing that sustained home prices at or above 5x median income consistently preceded displacement pressure. Research from the National Association of Realtors and the Federal Reserve Bank of Atlanta identifies price-to-income ratios above 5.0 as structural affordability stress indicators, particularly in metros with limited housing supply elasticity.

HPI as a Leading Indicator

The FHFA House Price Index (HPI) used extensively in this project is one of the most reliable and geographically consistent measures of residential price appreciation. Studies by Case-Shiller and the Federal Reserve have shown that sustained 3-year HPI growth above 25-30% signals market overheating. This project’s threshold of 30% over 3 years (`hpi_3yr_growth >= 0.30`) is grounded in this empirical literature.

Momentum-Based Detection

Shiller’s work on “irrational exuberance” and momentum pricing in real estate markets provides theoretical justification for the sustained-momentum condition used here. A single quarter of high YoY (year-over-year) growth is insufficient to identify structural change; consistent growth above 10% for 3+ quarters signals a regime shift, not a transient fluctuation.

3. Data Sources & Warehousing

All data is joined at the CBSA (Core-Based Statistical Area) level using CBSA codes at the primary key.

Source	Variable(s)	Granularity
Zillow Research (ZHVI)	Median home values	Monthly -> Quarterly
FHFA House Price Index	Price appreciation index	Quarterly
U.S. Census ACS (B19013_001)	Median household income	Annual
Census Population Estimates	Metro population	Annual

Tampa CBSA Note: Tampa’s CBSA code requires special handling. The official Census ACS GEOID for Tampa-ST.Petersburg-Clearwater metro is 45300, but the ZHVI data uses 45294 (Tampa MSAD). A crosswalk correction was applied in R:

```
None
income_clean_tampa_fix <- income %>%
  filter(cbsa == 45300) %>%
  mutate(cbsa = 45294) %>%
income_clean <- bind_rows(income_clean, income_clean_tampa_fix)
```

The merged dataset contained **83,699 rows x 25 columns** spanning quarterly observations across all U.S. CBSAs from 2010-2024. The three training cities (Austin: CBSA 12420, Boise: 14260, Tampa: 45294) are explicitly verified in the merged dataset before proceeding with feature engineering.

4. Affordability Collapse Definition & Threshold Selection

Before feature engineering, a precise, operationalizable definition of “affordability collapse onset” was required. Three conditions were tested and adopted:

Condition 1 - Price Acceleration

3-year HPI growth $\geq$ 30% (`hpi_3yr_growth >= 0.30`)

This captures rapid price escalation. Academic literature (Malpezzi & Wachter, 2005; Case & Shiller, 1989) documents that 3-year real HPI growth above 25-35% is strongly associated with market overheating and subsequent affordability deterioration.

Condition 2 - Sustained Momentum

YoY ZHVI growth $\geq$ 10% for 3+ consecutive quarters (zhvi_yoy $\geq$ 0.10)

A single-quarter spike may be noise; 3 consecutive quarters of $\geq$ 10% YoY ZHVI growth signals a durable trend. The 10% threshold aligns with Federal Reserve research characterizing “rapid” home price growth.

Condition 3 - Affordability Burden (Primary Label Criterion)

Median home price $>$ 5.0x median household income (price_to_income_ratio $\geq$ 5.0)

The PTI threshold of 5.0 was empirically selected after testing four alternatives (4.0, 4.5, 5.0, 5.5) against all three training cities. The threshold sensitivity analysis, conducted in Python, yielded the following onset dates:


Table 1: Affordability Collapse Onset Dates by PTI Threshold

Threshold	Austin Onset	Boise Onset	Tampa Onset
4.0	2014 Q3	2016 Q2	2017 Q4
4.5	2020 Q4	2017 Q4	2021 Q2
5.0	2021 Q2	2019 Q3	2021 Q4
5.5	2021 Q3	2020 Q4	2022 Q2

5. Exploratory Data Analysis

Figure 1: Price-to-Income Ratio Trajectories (2010-2024) - Austin, Boise, Tampa
The plot displays quarterly PTI for all three training cities from 2010 through 2024. A dashed horizontal line marks the 5.0 affordability collapse threshold.

Price-to-Income Ratio: Austin, Boise, Tampa



Key Findings: All three training cities crossed 5.0x PTI threshold during 2019-2021, with Boise the earliest and Austin/Tampa crossing nearly simultaneously during the COVID-era housing surge.

YoY ZHVI Growth Trends

YoY ZHVI growth for all three cities shows a pronounced acceleration beginning in 2020, consistent with COVID-19-era migration patterns, remote work adoption, and historically low mortgage rates driving demand well ahead of supply.

Population Velocity

Population velocity (year-over-year change in annual population growth) is available for Austin and Boise, confirming accelerating in-migration in the 2-3 years preceding PTI collapse. Tampa does not have population data in this pipeline. This is an acknowledged limitation (see Section 8).

EDA-To-Feature Engineering Bridge

The steep post-2020 acceleration visible in the YoY ZHVI growth trends above directly motivated the 3-quarter sustained-momentum condition (Condition 2) used in the collapse definition. A single-quarter spike, as seen sporadically in pre-2019 data for all three cities, does not appear in Figure 1's PTI trajectories as a lasting threshold crossing, whereas the sustained 2020-2021 climb does. This visual pattern, more than the numeric thresholds alone, justified excluding transient outliers from the `collapse_onset` label.

6. Feature Engineering

All features are lagged by 4 quarters (1 year) before any rolling computation. This is a critical data integrity decision: it ensures no feature accidentally incorporates same-quarter or future information that overlaps with the construction of the `collapse_onset` label, preventing data leakage.

Design Rule: *Never let a feature use current or future information that overlaps with label construction.*

Features Engineered

Feature 1: Price-to-Income 5-Year Change (`price_to_income_5yr_chg`)

The change in PTI from 20 quarters (5 years) prior, computed on the lagged PTI series. This captures structural, slow-moving affordability deterioration.

Python

```
df['price_to_income_lag'] = g['price_to_income_ratio'].shift(LAG)
df['price_to_income_5yr_chg'] = df.groupby('cbsa')['price_to_income_lag'].transform(lambda x: x - x.shift(20))
)
```

Feature 2: Affordability Momentum (`affordability_momentum`)

A 12-quarter rolling linear slope of lagged YoY ZHVI growth. This is a smooth, directional signal capturing whether price growth is accelerating or decelerating over a 3-year window.

Python

```
df['affordability_momentum'] = df.groupby('cbsa')['zhvi_yoy_lag'].transform(lambda x: x.rolling(12,
min_periods=12).apply(
    lambda y: np.polyfit(np.arange(len(y)), y, 1) [0] if y.notna().all() else np.n
)
)
```

Feature 3: HPI YoY Lag (`hpi_yoy_lag`)

Year-over-year change in the FHFA House Price Index, lagged 1 year. This is the most broadly available feature and the only one with multi-city holdout coverage (66,325 rows across 407 cities).

Feature 4&5: Population Velocity & Acceleration (`pop_velocity_lag`, `pop_acceleration_lag`)

First and second difference of annual population YoY growth, lagged 1 year. Available for Austin and Boise but not for **Tampa**. This is due to a known data gap. Used exclusively in model 2 (generalization/scoring model) and holdout scoring.

Table 2: Feature Coverage by City

Feature	Austin (non-null)	Boise (non-null)	Tampa (non-null)
affordability_momentum	85	81	145
zhvi_qoq_lag	99	95	159
price_to_income_5yr_chg	40	40	100
hpi_yoy_lag	187	179	244
pop_velocity_lag	55	55	0
pop_acceleration_lag	54	54	0

Tampa's population feature gap results in pop_velocity_lag = 0, requiring separate model architectures for explanatory vs. generalization use cases.

7. Predictive Modeling

Two models were developed with distinct purposes. Model 1 is designated as this project's primary model for answering the research question, since it uses the strongest available feature set and achieves the highest predictive accuracy. Model 2 exists specifically to extend scoring coverage to cities lacking Model 1's full feature set, and its lower performance is an accepted trade-off for that broader reach, not a competing claim to primacy.

Model 1 - Explanatory model (All 3 Training Cities, 4 Features)

Purpose: Identify housing indicators that are linked to current unaffordability and explain importance with SHAP. Because both the target and main feature use PTI data, Model 1 explains current unaffordability rather than forecasting future collapse.

Target Variable: is_unaffordable (binary: PTI>5.0)

Features: affordability_momentum, zhvi_qoq_lag, price_to_income_5yr_chg, hpi_yoy_lag

Training Cities: Austin, Boise, Tampa

Split: 70/30 stratified random split (train : {0: 85, 1:41}; val: {0: 37, 1:17})

Architecture: XGBoost classifier with shallow depth to prevent overfitting on the limited sample:

Python

```
model1_fresh = xgb.XGBClassifier( n_estimators=100, max_depth=3, learning_rate=0.05, subsample=0.8,
                                  colsample_bytree=0.8, eval_metric="logloss", random_state=42)
```

Table 3: Model 1 Performance (Validation Set)

Metric	Score
Accuracy	90.7%
Precision	83.3%
Recall	88.2%
F1-Score	85.7%
ROC-AUC	0.943

Classification Report:

Python				
	precision	recall	f1-score	support
0	0.94	0.92	0.93	37
1	0.83	0.88	0.86	17
accuracy			0.91	54

SHAP Feature Importance - Model 1

SHAP (SHapley Additive exPlanations) values were computed using shap.

Figure 2: SHAP Summary Plot - Model 1

This is a beeswarm plot that shows the distribution of SHAP values for each feature across all training observations. Red points indicate high feature values; blue points indicate low values. Rightward displacement indicates a positive impact on classifying a quarter as “unaffordable”.




Table 4: Mean Absolute SHAP Values - Model 1

Feature	Mean
price_to_income_5yr_chg	1.083
affordability_momentum	0.702
hpi_yoy_lag	0.588
zhvi_qoq_lag	0.138

Analysis: The 5-year change in price-to-income ratio is by far the dominant feature, contributing more than the other three features combined. This aligns with theory: slow-moving structural deterioration in affordability is a stronger collapse signal than short-term price momentum. Affordability momentum (the rolling slope of YoY ZHVI growth) ranks second, confirming that the *direction and acceleration* of price change matters beyond raw growth rates.

Cross-Table Synthesis:

Table 1's threshold sensitivity analysis and Table 4's SHAP importance rankings tell a consistent story: Boise's earlier onset across every PTI threshold (2016 Q2 at the 4.0x cutoff, versus 2017 Q4 for Austin and Tampa) corresponds with a steeper price_to_income_5yr_chg trajectory in the underlying feature data, reinforcing why this feature carries nearly double the mean absolute SHAP value of the second-ranked affordability_momentum (1.083 vs 0.702). This consistency across an independent sensitivity analysis (Table 1) and a model-internal interpretability method (Table 4) strengthens

confidence that the 5-year structural drift in affordability, not short-term price momentum, is the strongest feature associated with concurrent unaffordability in this training sample. Finding consistent with Malpezzi and Wachter’s (2005) framework distinguishing structural from speculative price deterioration.

Model 2 - Generalization/Scoring Model (Austin & Boise, 3 Features)

Purpose: Score 371 holdout metros area not in the training set.
Target Variable: is_unaffordable
Features: hpi_yoy_lag, pop_velocity_lag, pop_acceleration_lag
Training Cities: Austin and Boise
Split: 70/30 stratified random split (train : {0: 53, 1:22}; val: {0: 23, 1:10})

Hyperparameters were optimized using Optuna (100 trials, TPE sampler, 5-fold stratified CV, F1 objective), including *scale_pos_weight* tuning to account for class imbalance.

Table 5: Model 2 Performance (Validation Set)


Metric	Score
Accuracy	78.8%
Precision	71.4%
Recall	50.0%
F1-Score	58.8%
ROC-AUC	0.713

Model 2's reduced performance relative to Model 1 is expected and structurally appropriate. Model 2 uses only 3 broadly available features and is trained on only 2 cities- a deliberate trade-off to maximize holdout coverage (19,432 rows, 371 cities), at the cost of precision on the training cities themselves.

Class Imbalance in Results Interpretation: The class imbalance addressed via *scale_pos_weight* tuning during training carries a direct interpretive consequence. Model 2’s 50.0% recall means that roughly half of all quarters that were truly “unaffordable” in the Austin/Boise validation were missed entirely. This is not simply a method footnote, but it shows how much confidence to place in any single holdout metro’s *risk_score*, since the same recall ceiling likely applies when scoring the 371 unseen cities. A holdout metro genuinely on a collapse trajectory has a meaningfully higher chance of being under-flagged by Model 2 than by Model 1, reinforcing that Model 2’s risk scores should be read as a coarse screening signal, not a precise probability.

Figure 3: SHAP Summary Plot - Model 2

SHAP values were computed for Model 2 using the same Tree Explainer methodology applied to Model 1, revealing a markedly different importance structure than the explanatory model.



Analysis: Unlike Model 1, where `price_to_income_5yr_chg` dominated, Model 2's SHAP beeswarm plot shows `hpi_yoy` as the overwhelmingly dominant feature, with SHAP values running from approximately -2 to +2.5. A spread roughly 5-8x wider than either population feature. High feature values (pink/red points) for `hpi_yoy_lag` push strongly positive (toward classifying a metro as unaffordable), while low values (blue points) push strongly negative, indicating a clean, monotonic relationship between rapid home price appreciation and collapse risk.


By contrast, `pop_acceleration_lag` and `pop_velocity_lag` show much tighter SHAP value clustering near zero, with only modest spread (roughly -0.5 to +0.5 for acceleration, and even narrower band for velocity). This suggests that while population dynamics contribute some signal, they play a supporting rather than primary role in Model 2's classification decisions. This contrasts with the original hypothesis that population growth would be a strong independent driver, given its established role in the qualitative narrative around Austin's and Boise's in-migration-driven collapse.

This finding has an important implication for the Holdout City Risk Scoring discussed next. The `hpi_yoy_lag` is both the most broadly available feature (66,325 rows across 407 cities) and the dominant driver of Model 2's predictions; the 371-city risk scores are driven primarily by price appreciation momentum rather than population pressure. This helps explain why a small, high-amenity market like Bozeman, MT, has experienced extreme home price acceleration over the past couple of years, despite a comparatively smaller population base than Austin or Boise. This suggests the model can still rank at the top of the City Risk Index despite structural differences in growth patterns, since it weights price appreciation far more heavily than demographic velocity.

Holdout City Risk Scoring

Model 2 was applied to all 371 holdout metros to generate a risk_score (probability of being unaffordable), averaged across all available quarters per metro. The results are a ranked City Risk Index.

Figure 4: Holdout City Risk Scores



As shown in Figure 4, Bozeman, MT leads the top-15 ranked holdout metros by risk_score, followed by Pinehurst-Southern Pines, NC and Amherst Town-Northampton, MA. The visible drop-off in risk_score after the top few cities suggests a distinct cluster of highest-concern metros rather than a smooth continuum, meaning stakeholders should treat the top 3-5 ranked cities as materially more of a priority than the remainder of the top 15.

Note on label absence: Ground-truth collapse_onset labels exist only for Austin, Boise, and Tampa in this dataset. Holdout risk scores represent probabilistic early-warning signals, not validated predictions. This is consistent with the research framing: the goal is prospective flagging, not retrospective accuracy measurement.

Why Bozeman, MT Ranks First in the City Risk Index

Bozeman, Montana emerges as the highest-risk metro in the 371-city holdout scoring, and independent housing market data strongly corroborates this model output. As of early 2024, Bozeman's median home price approached \$1 million, and Gallatin County's 2023 median home price of roughly \$790,000 required an annual household income near \$181,000 to afford. This is nearly double the county's actual median income of approximately \$101,000, implying a price-to-income ratio of ~7.8x. This is well above the project's 5.0x collapse threshold (HCN, 2025; KBZK, 2026).

A 2025 economic report commissioned by the City of Bozeman found that homeowners need to earn 182% of area median income, or ~\$220,000 annually. To afford the median home independently, confirming the severity the model's risk_score is designed to detect, even without a verified collapse_onset label for Bozeman in the training data (Montana Free Press, 2026). Montana ranked as the least affordable state for homebuyers in a 2024 analysis, and by mid-2026 reporting, Montana's statewide home price-to-income ratio was above both California's and New York's (Bozeman Daily Chronicle, 2024).

What a False Positive or False Negative Means for Stakeholders

Because Model's 2 holdout scores cannot be validated against ground-truth labels, the practice stakes of misclassification deserve explicit treatment. A false-positive, meaning a metro incorrectly flagged as high-risk. This carried real economic cost as it could trigger premature rent-control advocacy, deter investment, or cause a regional lender to tighten mortgage underwriting in a market that was never actually on a collapse trajectory, potentially harming the very residents the model aims to protect. A false negative meaning a metro that is genuinely approaching collapse but scores low is worse than a False Positive. This means displacement pressure builds unflagged, and the early-warning system failed at the one job it was built for.

Given Model 2's validated recall of 50.0% (Table 5), stakeholders should treat a "low risk" holdout score with skepticism rather than as reassurance, particularly for metros with rapid recent price appreciation that may simply fall outside the model's detection window. This asymmetry argues for using the City Risk Index as a prioritization tool for further investigation, not as a stand-alone decision input for policy or lending.

8. Ethics & Limitations

Ethical Considerations

Data source bias: The model is trained exclusively on three Sun Belt metros that became unaffordable during the post-2019 housing surge. All three are relatively young, fast-growing cities with limited historical housing stock constraints. The model may systematically underweight affordability-collapse pathways that manifest differently in cities that collapsed before 2019.

Threshold subjectivity: The 5.0x PTI threshold, while empirically supported and consistent with existing housing literature, is still a human-defined cutoff. Some metros may experience functional displacement at 4.5x (e.g., metros with lower-income distributions), while others may remain stable at 5.5x. The threshold sensitivity analysis (Table 1) documents how results shift under alternative thresholds.

Proxy labeling: The project uses $PTI > 5.0$ as a proxy for affordability collapse. Actual collapse - displacement, community disruption, municipal budget strain is more complex and dimensional than any single numeric threshold. The label is a tractable operationalization, not a comprehensive definition.

A “high-risk” label can itself move markets: Because this project’s output is a ranked, publicly interpretable risk score, labeling a metro “high-risk” could become a self-fulfilling or self-defeating signal. Investors may treat the label as a buy signal ahead of anticipated appreciation, or it may discourage needed development in a misclassified market. This risk is inherent to a predictive tool that names specific places, and it argues for presenting the City Risk Index as a screening prioritization aid for further human review, not as a stand-alone public-facing ranking or automated decision tool.

Limitations:

Limitations	Description
Small training set	Only 3 cities have a ground-truth collapse label. Model 1 trains on ~126 rows.
Tampa population gap	Tampa lacks population data in the pipeline, preventing use of pop_velocity_lag in the full 3-city model.
No holdout labels	The 371 holdout metros have no verified collapse labels; risk scores cannot be externally validated against known outcomes.
Temporal coverage	ZHVI data begins ~2010; features requiring long lags (e.g., 5-year PTI change) are unavailable for the earliest quarters.
Macro confounders	The 2020-2022 housing surge was driven partly by COVID-specific factors (remote work, historically low mortgage rates) that may not recur. The model may over-fit to a specific macro regime.
CBSA mismatch	Tampa requires a manual CBSA crosswalk correction (45300 -> 45294), and any similar mismatches in holdout cities could introduce silent join failures.

Geographic training bias in holdout rankings	All three training cities are Sun Belt/mountain-west metros. Risk scores for structurally different regions. Specifically those in the Rust Belt and Northeast metros, which may face affordability stress through slow-growth pathways rather than rapid price appreciation. This should be treated with reduced confidence, since the model was never trained on that collapse pattern. Any Rust Belt or Northeast metro appearing in the top-15 holdout lists warrants additional qualitative scrutiny before being treated as comparably reliable to a Sun Belt or mountain-west match like Bozeman.
--	--

9. Conclusion

This project demonstrates that lagged, publicly available housing and economic indicators can distinguish quarters of concurrent unaffordability in the three training cities. Within this sample, the 5-year change in price-to-income ratio is the strongest feature associated with unaffordability. Model 1's 0.943 ROC-AUC reflects performance on a stratified random test set for concurrent classification, not validated future-collapse forecasting.

Directly answering this project's research question: *price-to-income 5-year change*, *affordability momentum*, and *HPI year-over-year growth* are the housing market indicators most strongly associated with concurrent unaffordability in this sample. Model 1 demonstrates these indicators and can identify at-risk cities with strong accuracy (90.7%) among cities structurally similar to the training set. Model 2 uses the same indicators in reduced form and can extend identification to a much broader set of metros. The downside is that confidence is scaled down for cities structurally dissimilar to Austin, Boise, and Tampa.

The two-model architecture- one for interpretability, one for broad scoring balances both the explanatory depth and geographic coverage. The holdout scoring pipeline applied to 371 metros produces a deployable City Risk Index that city planners, regional banks, and housing advocates can use as a forward-looking resource. The Bozeman, MT case study demonstrates this index's practical value. Despite lacking a ground-truth label, the model's top-ranked risk score aligns closely with independently reported, real-world affordability distress in that market. This suggests the City Risk index has genuine early-warning utility beyond the three training cities, provided any stakeholder using it keeps its false-negative risk and geographic training bias in view.

Future work should prioritize:

1. Expanding training labels to additional collapsed metros (e.g., Nashville, Phoenix, Denver, and at least one Rust Belt or Northeast metro to test cross-regional generalization).
2. Resolving Tampa's data gap
3. Incorporating permit issuance data as a supply-side feature
4. Building an R Shiny or Streamlit dashboard for interactive risk exploration

5. Reformulate the target to predict threshold crossing 1-2 years ahead using only information available before the prediction window.

References

Bozeman housing market cools, but a stable market hasn't meant an affordable one. (2026, April 2). *Montana Free Press.*

<https://montanafreepress.org/2026/04/03/bozeman-housing-market-cools-but-a-stable-market-hasnt-meant-an-affordable-one/>

Case, K. E., & Shiller, R. J. (1989). The efficiency of the market for single-family homes. *American Economic Review*, 79(1), 125–137.

Federal Housing Finance Agency. (2024). *House Price Index (HPI)*. FHFA.

<https://www.fhfa.gov/data/hpi>

Federal Reserve Bank of Atlanta. (n.d.). *Home Ownership Affordability Monitor*. Federal Reserve Bank of Atlanta.

<https://www.atlantafed.org/center-for-housing-and-policy/data-and-tools/home-ownership-affordability-monitor>

Let's break down the 2026 Gallatin Valley housing report. (2026, April 28). *KBZK.*

<https://www.kbzk.com/news/local-news/lets-break-down-the-2026-gallatin-valley-housing-report>

Malpezzi, S., & Wachter, S. M. (2005). The role of speculation in real estate cycles. *Journal of Real Estate Literature*, 13(2), 143–164.

Montana Governor's Housing Task Force. (2024). *Phase III final report*. Montana Department of Environmental Quality. <https://deq.mt.gov/files/About/Housing/HTF-PIII-final-report-2024.pdf>

Montana listed as least affordable US state for homebuyers. (2024, October 7). *Bozeman Daily Chronicle.*

https://www.bozemandailychronicle.com/news/state/montana-housing-crisis/article_dbd57449-6044-5121-9d9d-722e40ce146f.html

The toll of Bozeman's housing crisis. (2025, April 30). *High Country News.*

<https://www.hcn.org/issues/57-5/the-toll-of-bozemans-housing-crisis/>

U.S. Census Bureau. (2024). *American Community Survey, Table B19013: Median household income.*

U.S. Census Bureau. <https://data.census.gov/>

U.S. Department of Housing and Urban Development. (2025). *Comprehensive housing market analysis: Bozeman, Montana*. HUD User.

<https://www.huduser.gov/portal/publications/pdf/BozemanMT-CHMA-25.pdf>

U.S. Department of Housing and Urban Development. (n.d.). *Affordable housing*. HUD.gov.

https://www.hud.gov/topics/affordable_housing

Wachter, S. M. (2016). Credit supply and housing prices in national and local markets. NBER Working Paper Series.

Zillow Research. (2024). *Zillow Home Value Index (ZHVI) methodology*. Zillow.
<https://www.zillow.com/research/methodology-zhvi/>

Appendix

Appendix A: Key Variables Definitions

Variable	Definition
price_to_income_ratio	Quarterly median ZHVI / annual ACS median household income
collapse_onset	First quarter where PTI first exceeds 5.0 after a period below 5.0
is_unaffordable	Binary flag: $PTI > 5.0$
hpi_3yr_chg	$(HPI_t / HPI_{\{t-12\}}) - 1$; 3-year price appreciation
zhvi_yoy	$(ZHVI_t / ZHVI_{\{t-4\}}) - 1$; year-over-year home value growth
affordability_momentum	12-quarter rolling OLS slope of lagged zhvi_yoy
price_to_income_5yr_chg	$PTI_{lag} - PTI_{lag_20}$; 5-year structural drift in affordability
pop_velocity_lag	First difference of annual population YoY growth, lagged 4 quarters

Appendix B: Model Architecture Summary

	Model 1	Model 2
Purpose	Explanatory / SHAP	Generalization / Scoring
Training cities	Austin, Boise, Tampa	Austin, Boise
Features	4 (ZHVI + HPI + PTI)	3 (HPI + Population)
Train rows	126	75
Validation rows	54	33
ROC-AUC	0.943	0.713
Holdout cities	N/A	371